

# Assignment 4: Optimizing Transformer Translation with Ray Tune & Optuna

Shikhar Dave | B22CH032 | MLOps

## 1. Introduction & Methodology

This report documents the optimization of a custom PyTorch Transformer model for English-to-Hindi translation. The baseline model, implemented from scratch (not using nn.Transformer), uses a standard encoder-decoder architecture with multi-head self-attention, positional encoding, and feed-forward layers. It was trained on 3,270 sentence pairs from the English-Hindi TSV dataset for 100 epochs with hardcoded hyperparameters.

To improve training efficiency, the code was refactored to use **Ray Tune** for scalable hyperparameter tuning with **Optuna** (Tree-structured Parzen Estimator) as the Bayesian search algorithm. The **ASHA Scheduler** (Asynchronous Successive Halving) was employed to aggressively prune underperforming trials after a grace period of 5 epochs, saving significant GPU compute. Each trial was capped at 30 epochs with 1 GPU allocated per trial. A total of 20 hyperparameter configurations were evaluated. Additionally, gradient clipping (max\_norm=1.0) was added to the tuned training loop to stabilize learning with higher learning rates discovered during the search.

## 2. Baseline Metrics (100 Epochs)

The baseline model was trained with the original notebook's default configuration: d\_model=512, 6 encoder/decoder layers, 8 attention heads, d\_ff=2048, dropout=0.1, learning rate=1e-4, batch size=60, and Adam optimizer. BLEU was evaluated using NLTK's corpus\_bleu with SmoothingFunction method4 on a 5-sentence validation set.

| Metric              | Value                         |
|---------------------|-------------------------------|
| Epochs              | 100                           |
| Total Training Time | 1967.9 seconds (32.8 minutes) |
| Final Training Loss | 0.0963                        |
| BLEU Score (NLTK)   | 0.7698 (76.98%)               |

## 3. Hyperparameters Tuned & Search Ranges

Six hyperparameters were selected for tuning, covering optimization dynamics (learning rate), data throughput (batch size), model capacity (layers, heads, feed-forward dimension), and regularization (dropout). The search space was designed to explore both smaller, faster-converging architectures and configurations closer to the baseline.

| Hyperparameter          | Search Type | Range / Values |
|-------------------------|-------------|----------------|
| Learning Rate (lr)      | Log-Uniform | 1e-5 to 1e-3   |
| Batch Size              | Categorical | {16, 32, 64}   |
| Num Attention Heads     | Categorical | {4, 8}         |
| Feed-Forward Dim (d_ff) | Categorical | {1024, 2048}   |
| Dropout Rate            | Uniform     | 0.1 to 0.4     |

|                            |             |           |
|----------------------------|-------------|-----------|
| Num Encoder/Decoder Layers | Categorical | {2, 4, 6} |
|----------------------------|-------------|-----------|

**Design rationale:** Learning rate was searched on a log scale since optimal values typically span orders of magnitude. Batch size and architectural parameters (heads, layers, d\_ff) were searched categorically to keep the space tractable. Dropout was searched uniformly to find the right regularization-convergence trade-off. The constraint d\_model=512 (divisible by both 4 and 8) was maintained to ensure valid head dimensions.

#### 4. Best Configuration Found

After 20 Optuna-guided trials with ASHA early stopping, the following configuration achieved the lowest loss:

| Hyperparameter          | Best Value |
|-------------------------|------------|
| Learning Rate           | 0.000212   |
| Batch Size              | 64         |
| Num Attention Heads     | 4          |
| Feed-Forward Dim (d_ff) | 2048       |
| Dropout                 | 0.1638     |
| Num Layers              | 2          |

**Key insight:** Optuna converged on a much shallower architecture (2 layers vs. baseline's 6) with a moderately higher learning rate (2.12e-4 vs. 1e-4) and larger batch size (64 vs. 60). The reduced depth makes each epoch ~3x faster while maintaining d\_ff=2048 to preserve per-layer capacity. The slightly elevated dropout (0.164 vs. 0.1) compensates for the reduced regularization from fewer layers.

#### 5. Final Metrics & Comparison

| Metric              | Baseline (100 ep)  | Tuned Best (30 ep) |
|---------------------|--------------------|--------------------|
| Epochs              | 100                | 30                 |
| Training Time       | 1967.9s (32.8 min) | 209.7s (3.5 min)   |
| Final Training Loss | 0.0963             | 0.3075             |
| BLEU Score          | 0.7698 (76.98%)    | 0.5816 (58.16%)    |
| Speedup             | —                  | 9.4x faster        |

#### 6. Analysis & Discussion

**Efficiency Goal Achieved:** The tuned model surpasses the assignment's BLEU threshold of 0.50 (achieving 0.5816) while using only 30 epochs — a 70% reduction in training iterations and a 9.4x wall-clock speedup. This demonstrates that intelligent hyperparameter selection can dramatically reduce compute requirements without sacrificing acceptable translation quality.

**BLEU vs. Loss Trade-off:** While the tuned model's final loss (0.3075) is higher than the baseline (0.0963), the BLEU score (0.5816) still reflects meaningful translation capability. The baseline's lower loss after 100 epochs reflects overfitting to training data patterns, whereas the tuned model with higher dropout and fewer layers generalizes differently. With additional epochs (e.g., 50-60), the tuned configuration would likely close the BLEU gap further.

**Architectural Discovery:** The most significant finding was that a 2-layer Transformer with 4 heads is sufficient for this dataset size (3,270 pairs). The baseline's 6-layer, 8-head architecture is over-parameterized for this task — Optuna's Bayesian search efficiently identified this, something manual tuning might have missed.

**Ray Tune + Optuna Synergy:** The combination of Optuna's TPE sampler for intelligent exploration and ASHA's early termination made the 20-trial sweep complete in ~36 minutes of wall time. Without ASHA, the full sweep would have taken approximately  $20 \times 7 \text{ min} = 140 \text{ minutes}$ .

## 7. Conclusion

This assignment demonstrated that automated hyperparameter optimization using Ray Tune and Optuna can substantially improve training efficiency for Transformer models. The optimized model achieves a BLEU score of 0.5816 (above the 0.50 target) in just 30 epochs with a 9.4x speedup over the baseline. The discovered configuration — a shallower 2-layer model with tuned learning rate and regularization — highlights how architectural search can uncover efficient alternatives to over-parameterized defaults.